

```
In [1]: ## Import basic python libraries
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: # Import dataset
data1=pd.read_csv(r'C:\Users\anith\OneDrive\Documents\anil.csv')
```

```
In [3]: data1.head()
```

```
Out[3]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Terr
0	LP001002	Male	No	0	Graduate	No	5849	0.0	NaN	360.
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128.0	360.
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66.0	360.
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120.0	360.
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141.0	360.

```
In [4]: # Dealing the missing values
data1.isnull().sum()
```

```
Out[4]:
```

Loan_ID	0
Gender	13
Married	3
Dependents	15
Education	0
Self_Employed	32
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	22
Loan_Amount_Term	14
Credit_History	50
Property_Area	0
Loan_Status	0

dtype: int64

```
In [5]: # Remove the missing values
data1.dropna(inplace=True)
data1.isnull().sum()
```

```
Out[5]:
```

Loan_ID	0
Gender	0
Married	0
Dependents	0
Education	0
Self_Employed	0
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	0
Loan_Amount_Term	0
Credit_History	0
Property_Area	0
Loan_Status	0

dtype: int64

```
In [6]: data1.shape
```

```
Out[6]: (480, 13)
```

```
In [7]: data1
```

Out[7]:	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_T
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128.0	36
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66.0	36
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120.0	36
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141.0	36
5	LP001011	Male	Yes	2	Graduate	Yes	5417	4196.0	267.0	36
...
609	LP002978	Female	No	0	Graduate	No	2900	0.0	71.0	36
610	LP002979	Male	Yes	3+	Graduate	No	4106	0.0	40.0	18
611	LP002983	Male	Yes	1	Graduate	No	8072	240.0	253.0	36
612	LP002984	Male	Yes	2	Graduate	No	7583	0.0	187.0	36
613	LP002990	Female	No	0	Graduate	Yes	4583	0.0	133.0	36

480 rows × 13 columns

```

In [8]: ### Visualization###
def bar_chart(col):
    Approved = data1[data1["Loan_Status"]=="Y"][col].value_counts()
    Disapproved = data1[data1["Loan_Status"]=="N"][col].value_counts()

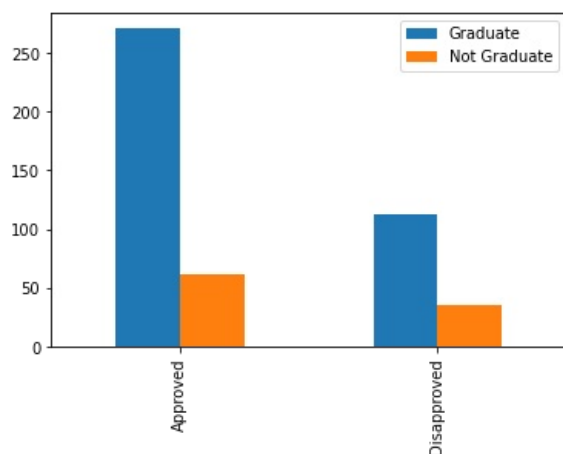
    df1 = pd.DataFrame([Approved, Disapproved])
    df1.index = ["Approved", "Disapproved"]
    df1.plot(kind="bar")

```

```

In [9]: bar_chart('Education')


```



```

In [10]: bar_chart('Gender')


```



```

In [11]: bar_chart('Credit_History')

```



```
In [12]: #checking the skewness (acceptable range is -5 to +5)
data1.skew()
```

```
Out[12]: ApplicantIncome      6.917027
CoapplicantIncome      5.881622
LoanAmount             2.361437
Loan_Amount_Term      -2.333710
Credit_History        -2.013253
dtype: float64
```

```
In [13]: data=data1.drop(['Loan_ID'],axis=1)
```

```
In [14]: data
```

```
Out[14]:
```

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit
1	Male	Yes	1	Graduate	No	4583	1508.0	128.0	360.0	
2	Male	Yes	0	Graduate	Yes	3000	0.0	66.0	360.0	
3	Male	Yes	0	Not Graduate	No	2583	2358.0	120.0	360.0	
4	Male	No	0	Graduate	No	6000	0.0	141.0	360.0	
5	Male	Yes	2	Graduate	Yes	5417	4196.0	267.0	360.0	
...
609	Female	No	0	Graduate	No	2900	0.0	71.0	360.0	
610	Male	Yes	3+	Graduate	No	4106	0.0	40.0	180.0	
611	Male	Yes	1	Graduate	No	8072	240.0	253.0	360.0	
612	Male	Yes	2	Graduate	No	7583	0.0	187.0	360.0	
613	Female	No	0	Graduate	Yes	4583	0.0	133.0	360.0	

480 rows × 12 columns

```
In [15]: data['Dependents'].unique()
```

```
Out[15]: array(['1', '0', '2', '3+', dtype=object)
```

```
In [16]: data['Dependents']=data['Dependents'].replace('3+',4)
```

```
In [17]: data['Dependents'].unique()
```

```
Out[17]: array(['1', '0', '2', 4], dtype=object)
```

```
In [18]: data['Dependents']=data['Dependents'].astype('int')
```

```
In [19]: data['Dependents'].unique()
```

```
Out[19]: array([1, 0, 2, 4])
```

```
In [20]: #by using ordinal encoder converting vectors
```

```
from sklearn.preprocessing import OrdinalEncoder
```

```
encoder = OrdinalEncoder()
```

```
data[["Gender","Married","Education","Self_Employed","Property_Area","Loan_Status"]] = encoder.fit_transform(data.head())
```

Out[20]:	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_H
1	1.0	1.0	1	0.0	0.0	4583	1508.0	128.0	360.0	
2	1.0	1.0	0	0.0	1.0	3000	0.0	66.0	360.0	
3	1.0	1.0	0	1.0	0.0	2583	2358.0	120.0	360.0	
4	1.0	0.0	0	0.0	0.0	6000	0.0	141.0	360.0	
5	1.0	1.0	2	0.0	1.0	5417	4196.0	267.0	360.0	

```
In [21]: #encoding the features are float we convert into integer
data[["Gender", 'Married', 'Education', 'Self_Employed', 'Property_Area', 'Loan_Status']] = data[["Gender", 'Married', 'Property_Area', 'Loan_Status']].astype(int)
```

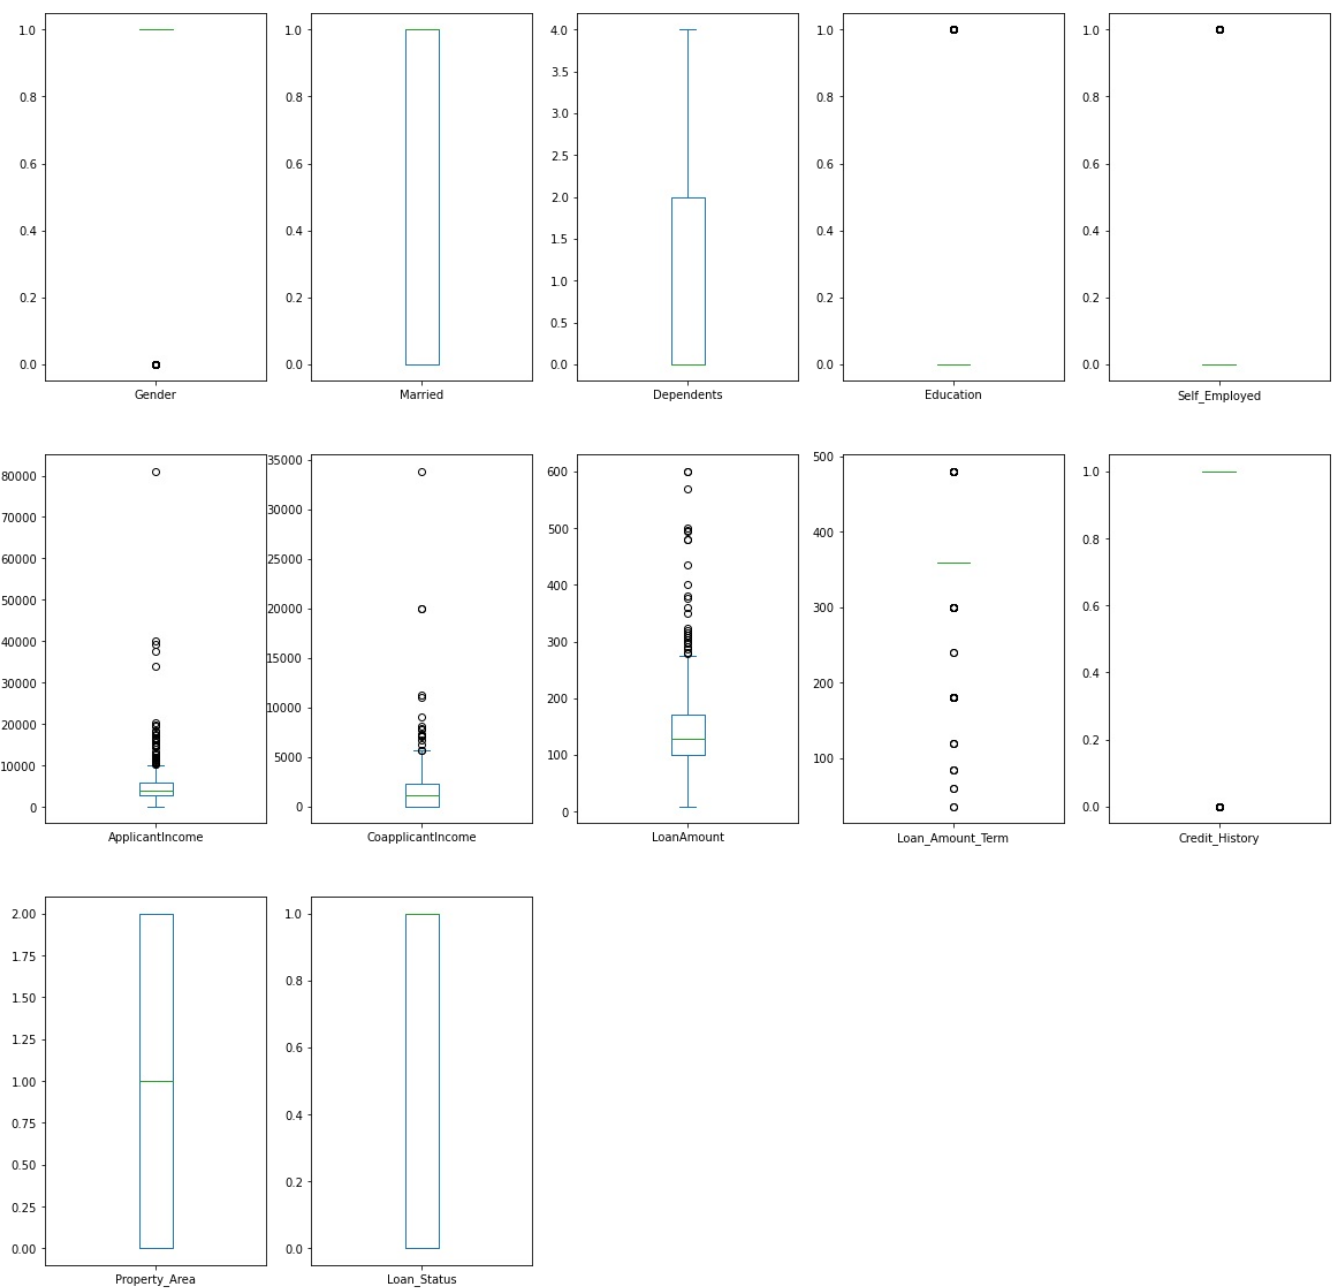
```
In [22]: data
```

Out[22]:	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_H
1	1	1	1	0	0	4583	1508.0	128.0	360.0	
2	1	1	0	0	1	3000	0.0	66.0	360.0	
3	1	1	0	1	0	2583	2358.0	120.0	360.0	
4	1	0	0	0	0	6000	0.0	141.0	360.0	
5	1	1	2	0	1	5417	4196.0	267.0	360.0	
...
609	0	0	0	0	0	2900	0.0	71.0	360.0	
610	1	1	4	0	0	4106	0.0	40.0	180.0	
611	1	1	1	0	0	8072	240.0	253.0	360.0	
612	1	1	2	0	0	7583	0.0	187.0	360.0	
613	0	0	0	0	1	4583	0.0	133.0	360.0	

480 rows × 12 columns


```
In [23]: #seeing the outliers
data.plot(kind='box',subplots=True,layout=(3,5),figsize=(20,20))
```

```
Out[23]: Gender      AxesSubplot(0.125,0.657941;0.133621x0.222059)
Married    AxesSubplot(0.285345,0.657941;0.133621x0.222059)
Dependents AxesSubplot(0.44569,0.657941;0.133621x0.222059)
Education  AxesSubplot(0.606034,0.657941;0.133621x0.222059)
Self_Employed AxesSubplot(0.766379,0.657941;0.133621x0.222059)
ApplicantIncome AxesSubplot(0.125,0.391471;0.133621x0.222059)
CoapplicantIncome AxesSubplot(0.285345,0.391471;0.133621x0.222059)
LoanAmount  AxesSubplot(0.44569,0.391471;0.133621x0.222059)
Loan_Amount_Term AxesSubplot(0.606034,0.391471;0.133621x0.222059)
Credit_History AxesSubplot(0.766379,0.391471;0.133621x0.222059)
Property_Area AxesSubplot(0.125,0.125;0.133621x0.222059)
Loan_Status AxesSubplot(0.285345,0.125;0.133621x0.222059)
dtype: object
```



```
In [24]: #checking correlation of dataset
plt.figure(figsize=(20,20))
sns.heatmap(data.corr(),annot=True)
```

Out[24]: <AxesSubplot:>



```
In [25]: #removing multicollinearity by using vif acceptable range (-10 to 10)
from statsmodels.stats.outliers_influence import variance_inflation_factor
vif=pd.DataFrame()
vif['features']=data.columns
vif['vif factor']=[variance_inflation_factor(data.values,i) for i in range(data.shape[1])]
vif
```

Out[25]:

	features	vif factor
0	Gender	6.146708
1	Married	3.716146
2	Dependents	1.789241
3	Education	1.283195
4	Self_Employed	1.195656
5	ApplicantIncome	2.752965
6	CoapplicantIncome	1.554019
7	LoanAmount	6.412868
8	Loan_Amount_Term	10.747859
9	Credit_History	8.934985
10	Property_Area	2.610845
11	Loan_Status	4.633473

In [26]: data=data.drop(['Loan_Amount_Term'],axis=1)
data

Out[26]:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Credit_History	Property_Area
1	1	1	1	1	0	0	4583	1508.0	128.0	1.0
2	1	1	0	0	0	1	3000	0.0	66.0	1.0
3	1	1	0	1	0	2583	2358.0	120.0	1.0	1.0
4	1	0	0	0	0	6000	0.0	141.0	1.0	1.0
5	1	1	2	0	1	5417	4196.0	267.0	1.0	1.0
...
609	0	0	0	0	0	2900	0.0	71.0	1.0	1.0
610	1	1	4	0	0	4106	0.0	40.0	1.0	1.0
611	1	1	1	0	0	8072	240.0	253.0	1.0	1.0
612	1	1	2	0	0	7583	0.0	187.0	1.0	1.0
613	0	0	0	0	1	4583	0.0	133.0	0.0	0.0

480 rows × 11 columns

In [27]: x=data.drop(['Loan_Status'],axis=1)
y=data['Loan_Status']

In [28]: # feature scaling
from sklearn.preprocessing import StandardScaler
std_scaler=StandardScaler()
x_scaled=std_scaler.fit_transform(x)

In [29]: x_scaled

Out[29]: array([[0.46719815, 0.73716237, 0.11235219, ..., -0.20808917,
 0.41319694, -1.31886834],
 [0.46719815, 0.73716237, -0.70475462, ..., -0.97900085,
 0.41319694, 1.25977445],
 [0.46719815, 0.73716237, -0.70475462, ..., -0.30756164,
 0.41319694, 1.25977445],
 ...,
 [0.46719815, 0.73716237, 0.11235219, ..., 1.34616826,
 0.41319694, 1.25977445],
 [0.46719815, 0.73716237, 0.92945899, ..., 0.52552034,
 0.41319694, 1.25977445],
 [-2.14041943, -1.35655324, -0.70475462, ..., -0.14591887,
 -2.42015348, -0.02954695]])

In [30]: # training the model
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x_scaled,y,test_size=0.3,random_state=45)

In [31]: x_train.shape

Out[31]: (336, 10)

In [32]: y_train.shape

Out[32]: (336,)

In [33]: #import varies models and metrics

```

from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, f1_score

```

```

In [34]: #testing of accuracy of models
models=[LogisticRegression(),SVC(),RandomForestClassifier(), DecisionTreeClassifier()]
for m in models:
    print(f'{m}:')
    m.fit(x_train,y_train)
    print('Training score:',m.score(x_train,y_train))
    print('Testing score:',m.score(x_test,y_test))
    predm=m.predict(x_test)
    flscore=f1_score(y_test,predm)
    print('flscore:',flscore)
    acrscore=accuracy_score(y_test,predm)
    print('Accuracy score:',acrscore)
    crsv=cross_val_score(m,x_scaled,y,cv=5)
    print('Cross validation mean score:',crsv.mean())
    print("")
    print('***5)
    print('\n')

```

```

LogisticRegression():
Training score: 0.8154761904761905
Testing score: 0.7916666666666666
flscore: 0.8648648648648648
Accuracy score: 0.7916666666666666
Cross validation mean score: 0.8020833333333334

```

```

SVC():
Training score: 0.8392857142857143
Testing score: 0.7847222222222222
flscore: 0.8597285067873304
Accuracy score: 0.7847222222222222
Cross validation mean score: 0.8104166666666666

```

```

RandomForestClassifier():
Training score: 1.0
Testing score: 0.7638888888888888
flscore: 0.8440366972477066
Accuracy score: 0.7638888888888888
Cross validation mean score: 0.7916666666666667

```

```

DecisionTreeClassifier():
Training score: 1.0
Testing score: 0.7361111111111112
flscore: 0.8190476190476189
Accuracy score: 0.7361111111111112
Cross validation mean score: 0.7104166666666666

```

by using hyper parameter tuning increase the select model accuracy

```

In [35]: lr=LogisticRegression()

```

```

In [36]: param_grid=[{'penalty':['l1','l2','elasticnet','none'],'C':np.logspace(-4,4,20),'solver':['lbfgs','newton-cg'],'

```

```

In [37]: from sklearn.model_selection import GridSearchCV

```

```

In [38]: cif=GridSearchCV(lr,param_grid=param_grid,cv=3,verbose=True,n_jobs=-1)

```

```

In [39]: best_cif=cif.fit(x,y)

```

Fitting 3 folds for each of 1600 candidates, totalling 4800 fits

```

In [40]: best_cif.best_estimator_

```



```
Out[40]: LogisticRegression
LogisticRegression(C=0.08858667904100823, penalty='l1', solver='liblinear')
```

```
In [41]: print('Best Accuracy score:',best_cif.score(x,y))
```

Best Accuracy score: 0.80625

```
In [42]: #dumping the model
import joblib
import pickle
```

```
In [43]: joblib.dump(best_cif,'model.pkl')
model=joblib.load('model.pkl')
model.predict(x_test)
```

```
Out[43]: array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1,
1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 0, 1, 1, 1, 0, 1, 1, 0, 1,
1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1])
```

```
In [ ]:
```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js